

Benchmark on Deep Learning Frameworks

김 형 준

ARTIFICIAL INTELLIGENCE

Early artificial intelligence stirs excitement.



MACHINE LEARNING

Machine learning begins to flourish.



DEEP LEARNING

Deep learning breakthroughs drive AI boom.



1950's

1960's

1970's

1980's

1990's

2000's


2010's

DEEP LEARNING

TRAINING


Learning a new capability
from existing data

Untrained
Neural Network
Model



Deep Learning
Framework

TRAINING
DATASET



Trained Model
New Capability




INFERENCE

Applying this capability
to new data

NEW
DATA



App or Service
Featuring Capability



Trained Model
Optimized for
Performance

imperative
symbolic

theano


Chainer

PYTORCH

mxnet

Caffe

K

Caffe2

 Microsoft
CNTK


TensorFlow

before

2012

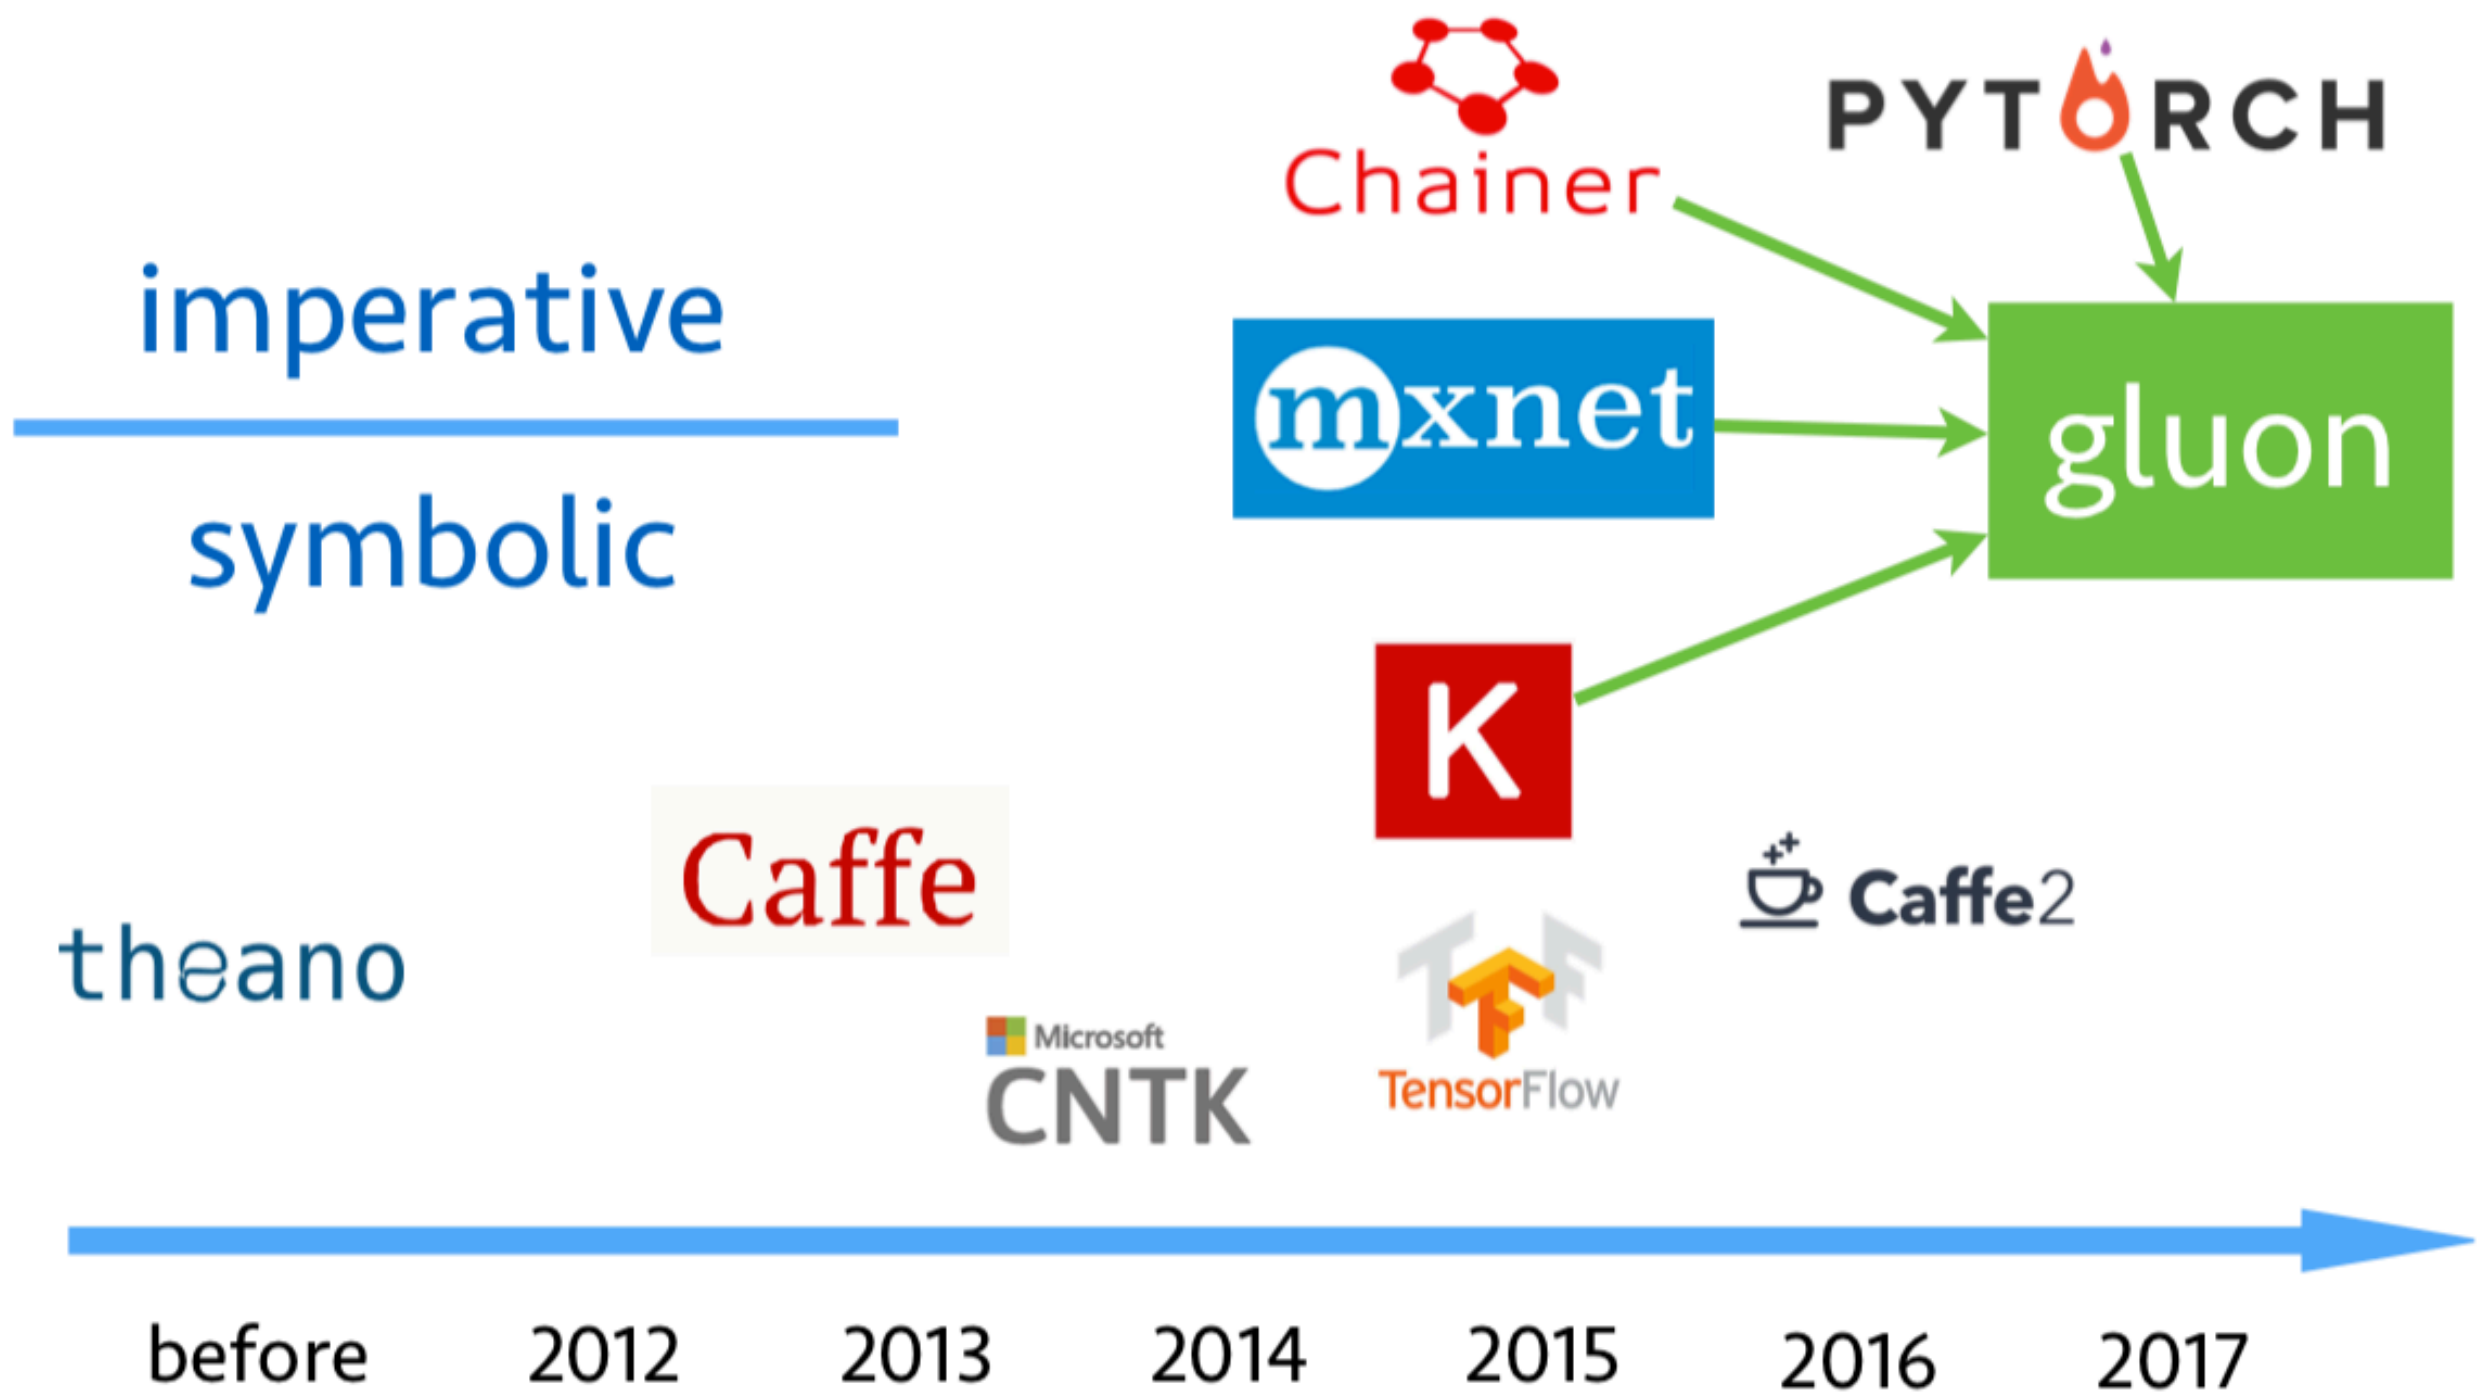
2013

2014

2015

2016

2017



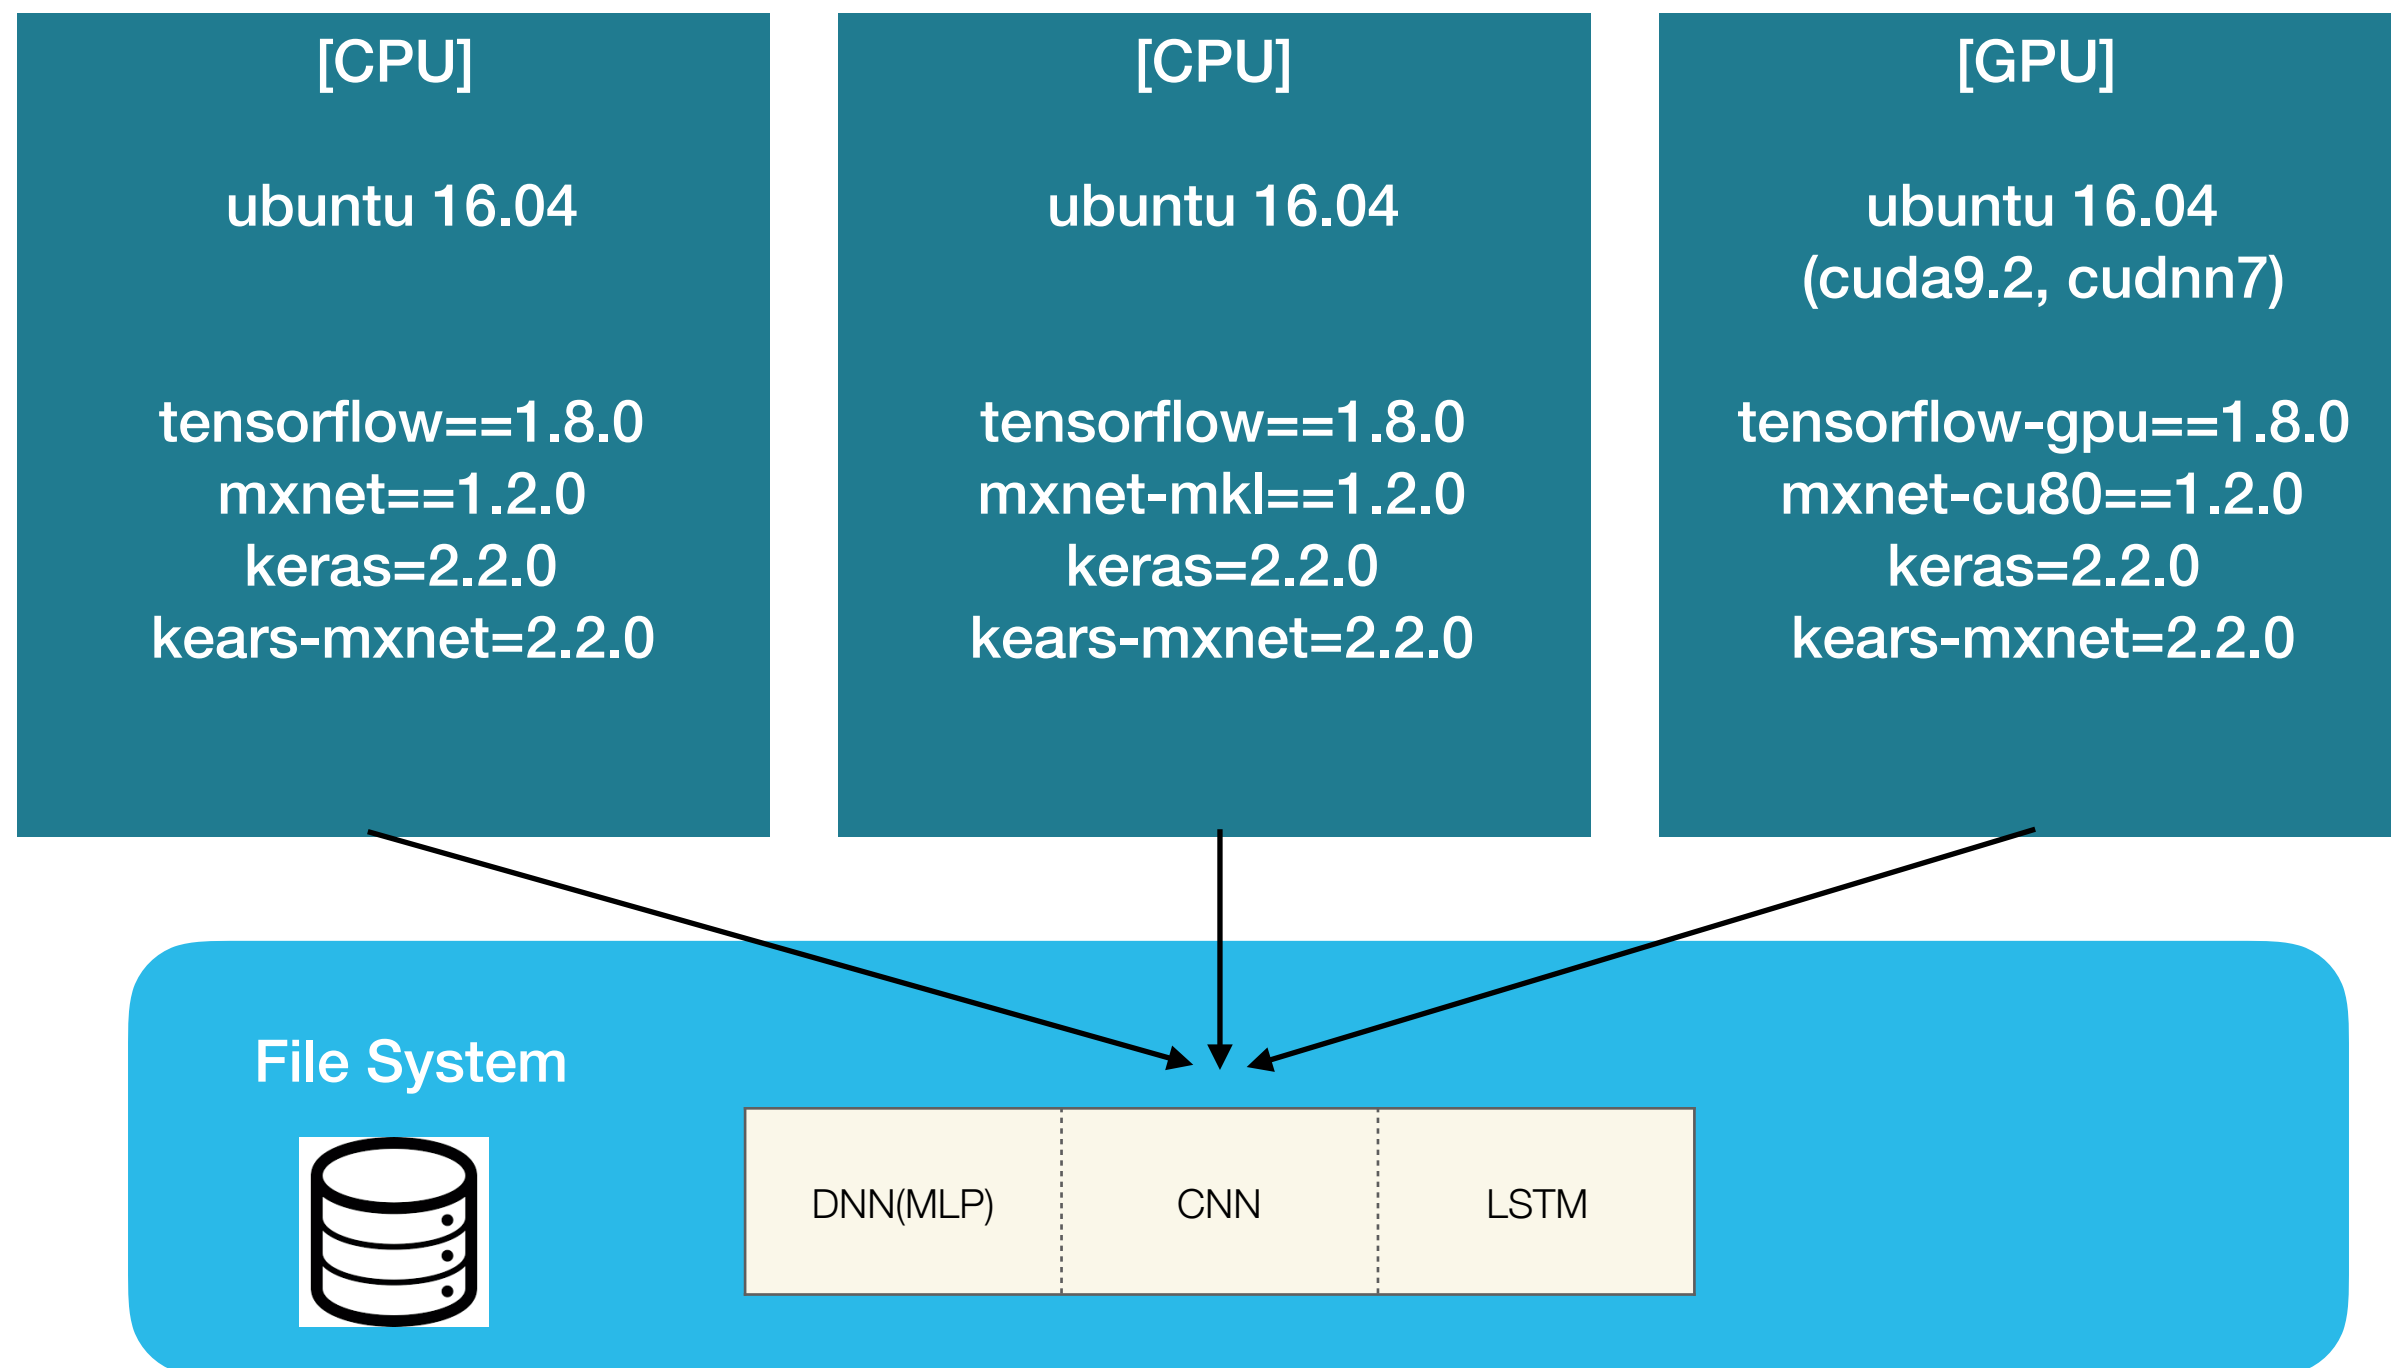

Framework Benchmark

https://github.com/tensorflow/benchmarks/tree/master/scripts/keras_benchmarks

https://github.com/awslabs/keras-apache-mxnet/tree/keras2_mxnet_backend/benchmark

Tensorflow		CPU (FULL-CORE)	GPU (1-GPU)
D	MXNET	CPU (FULL-CORE)	GPU (1-GPU)
	DNN(MLP)		
	CNN		
	LSTM		
	RN		


keras-mxnet



Between Platform - CPU


(ms)

DNN TRAIN



(ms)

DNN INFER



Tensorflow - CPU

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
47					@profile
48					def run_benchmark_train(self, gpus=0):
49					
50	1	2.0	2.0	0.0	x_train, y_train = self.x_train, self.y_train
51					
52	1	5569.0	5569.0	0.4	model = Sequential()
53	1	15978.0	15978.0	1.1	model.add(Dense(512, activation='relu', input_shape=(784,)))
54	1	19014.0	19014.0	1.3	model.add(Dropout(0.2))
55	1	14490.0	14490.0	1.0	model.add(Dense(512, activation='relu'))
56	1	17554.0	17554.0	1.2	model.add(Dropout(0.2))
57	1	14042.0	14042.0	1.0	model.add(Dense(self.num_classes))
58					
59	1	3.0	3.0	0.0	if keras.backend.backend() is "tensorflow" and gpus > 1:
60					model = multi_gpu_model(model, gpus=gpus)
61					
62	1	0.0	0.0	0.0	model.compile(loss='categorical_crossentropy',
63	1	10901.0	10901.0	0.8	optimizer=RMSprop(),
64	1	28131.0	28131.0	1.9	metrics=['accuracy'])
65					
66					# create a distributed trainer for cntk
67	1	2.0	2.0	0.0	if keras.backend.backend() is "cntk" and gpus > 1:
68					start, end = cntk_gpu_mode_config(model, x_train.shape[0])
69					x_train = x_train[start: end]
70					y_train = y_train[start: end]
71					
72	1	3.0	3.0	0.0	time_callback = timehistory.TimeHistory()
73					
74	1	2.0	2.0	0.0	model.fit(x_train, y_train, batch_size=self.batch_size,
75	1	1.0	1.0	0.0	epochs=self.epochs, verbose=1,
76	1	1326038.0	1326038.0	91.3	callbacks=[time_callback])
77					
78	1	5.0	5.0	0.0	self.fit_time = 0
79	2	3.0	1.5	0.0	for i in range(1, self.epochs):
80	1	1.0	1.0	0.0	self.fit_time += time_callback.times[i]
81					
82	1	0.0	0.0	0.0	self.model = model

mxnet-mkl - CPU


Line #	Hits	Time	Per Hit	% Time	Line Contents
47					@profile
48					def run_benchmark_train(self, gpus=0):
49					
50	1	3.0	3.0	0.0	x_train, y_train = self.x_train, self.y_train
51					
52	1	8234.0	8234.0	1.2	model = Sequential()
53	1	10504.0	10504.0	1.5	model.add(Dense(512, activation='relu', input_shape=(784,)))
54	1	565.0	565.0	0.1	model.add(Dropout(0.2))
55	1	6512.0	6512.0	1.0	model.add(Dense(512, activation='relu'))
56	1	706.0	706.0	0.1	model.add(Dropout(0.2))
57	1	5017.0	5017.0	0.7	model.add(Dense(self.num_classes))
58					
59	1	7.0	7.0	0.0	if keras.backend.backend() is "tensorflow" and gpus > 1:
60					model = multi_gpu_model(model, gpus=gpus)
61					
62	1	1.0	1.0	0.0	model.compile(loss='categorical_crossentropy',
63	1	368.0	368.0	0.1	optimizer=RMSprop(),
64	1	4409.0	4409.0	0.6	metrics=['accuracy'])
65					
66					# create a distributed trainer for cntk
67	1	1.0	1.0	0.0	if keras.backend.backend() is "cntk" and gpus > 1:
68					start, end = cntk_gpu_mode_config(model, x_train.shape[0])
69					x_train = x_train[start: end]
70					y_train = y_train[start: end]
71					
72	1	7.0	7.0	0.0	time_callback = timehistory.TimeHistory()
73					
74	1	2.0	2.0	0.0	model.fit(x_train, y_train, batch_size=self.batch_size,
75	1	0.0	0.0	0.0	epochs=self.epochs, verbose=1,
76	1	643038.0	643038.0	94.7	callbacks=[time_callback])
77					
78	1	4.0	4.0	0.0	self.fit_time = 0
79	2	4.0	2.0	0.0	for i in range(1, self.epochs):
80	1	2.0	2.0	0.0	self.fit_time += time_callback.times[i]
81					
82	1	1.0	1.0	0.0	self.model = model

mxnet-mkl - CPU


```
1  [|||||92.2%]
2  [|||||100.0%]
3  [|||||97.4%]
4  [|||||90.3%]
5  [|||||94.8%]
6  [|||||97.4%]
7  [|||||96.1%]
8  [|||||100.0%]
9  [|||||93.5%]
10 [|||||100.0%]
11 [|||||100.0%]
12 [|||||98.1%]
13 [|||||99.4%]
14 [|||||98.1%]
15 [|||||100.0%]
16 [|||||100.0%]
17 [|||||100.0%]
18 [|||||100.0%]
19 [|||||100.0%]
20 [|||||100.0%]
21 [|||||100.0%]
22 [|||||100.0%]
23 [|||||100.0%]
24 [|||||100.0%]
25 [|||||100.0%]
26 [|||||100.0%]
27 [|||||100.0%]
28 [|||||100.0%]
29 [|||||100.0%]
30 [|||||94.2%]
31 [|||||100.0%]
32 [|||||98.1%]
33 [|||||100.0%]
34 [|||||100.0%]
35 [|||||66.2%]
36 [|||||100.0%]
37 [|||||100.0%]
38 [|||||94.2%]
39 [|||||98.1%]
40 [|||||100.0%]
41 [|||||100.0%]
42 [|||||99.4%]
43 [|||||100.0%]
44 [|||||100.0%]
45 [|||||100.0%]
46 [|||||100.0%]
47 [|||||100.0%]
48 [|||||100.0%]
49 [|||||100.0%]
50 [|||||100.0%]
51 [|||||100.0%]
52 [|||||100.0%]
53 [|||||99.4%]
54 [|||||100.0%]
55 [|||||100.0%]
56 [|||||98.7%]
```

Mem [|||||5.846/7566] Task: 462 - 1407 threads running


CNN_TRAIN




CNN_INFER



LSTM_TRAIN




LSTM_INFER

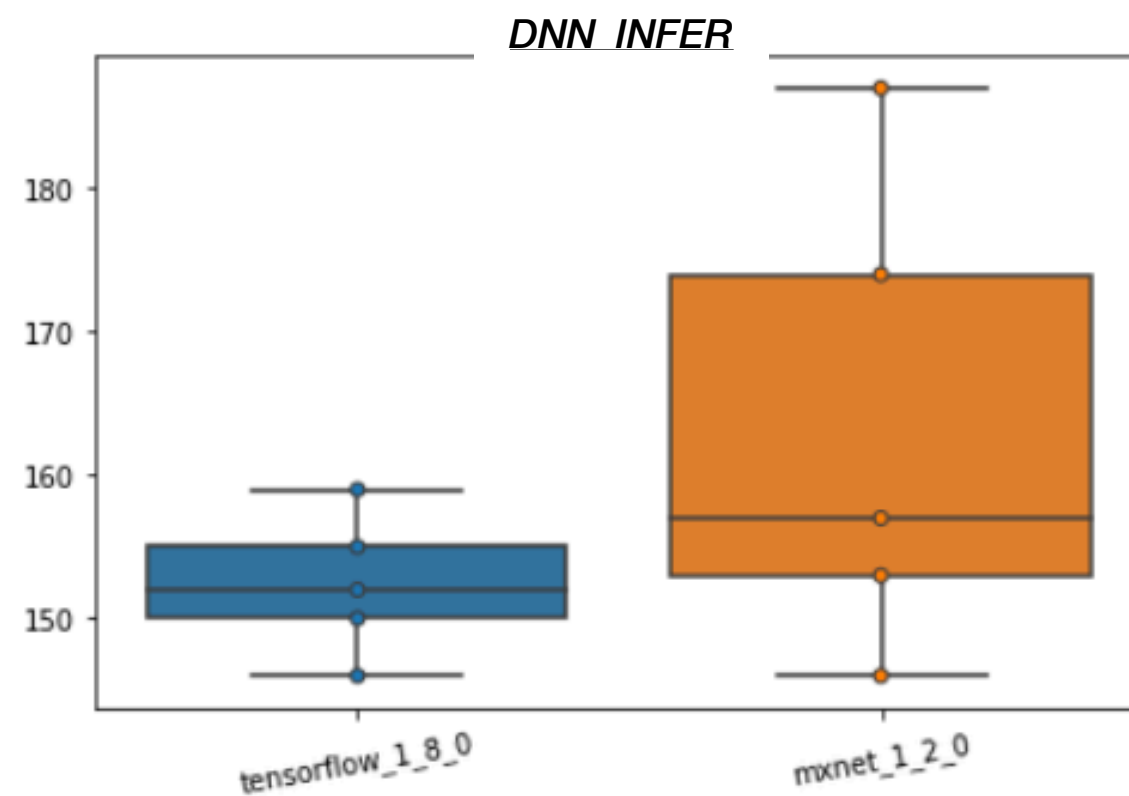


Between Platform - GPU

(ms)




(ms)




Between Platform - GPU


CNN TRAIN




CNN INFER



LSTM TRAIN



LSTM INFER



Tensorflow - CPU

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
47					@profile
48					def run_benchmark_train(self, gpus=0):
49					
50	1	2.0	2.0	0.0	x_train, y_train = self.x_train, self.y_train
51					
52	1	5569.0	5569.0	0.4	model = Sequential()
53	1	15978.0	15978.0	1.1	model.add(Dense(512, activation='relu', input_shape=(784,)))
54	1	19014.0	19014.0	1.3	model.add(Dropout(0.2))
55	1	14490.0	14490.0	1.0	model.add(Dense(512, activation='relu'))
56	1	17554.0	17554.0	1.2	model.add(Dropout(0.2))
57	1	14042.0	14042.0	1.0	model.add(Dense(self.num_classes))
58					
59	1	3.0	3.0	0.0	if keras.backend.backend() is "tensorflow" and gpus > 1:
60					model = multi_gpu_model(model, gpus=gpus)
61					
62	1	0.0	0.0	0.0	model.compile(loss='categorical_crossentropy',
63	1	10901.0	10901.0	0.8	optimizer=RMSprop(),
64	1	28131.0	28131.0	1.9	metrics=['accuracy'])
65					
66					# create a distributed trainer for cntk
67	1	2.0	2.0	0.0	if keras.backend.backend() is "cntk" and gpus > 1:
68					start, end = cntk_gpu_mode_config(model, x_train.shape[0])
69					x_train = x_train[start: end]
70					y_train = y_train[start: end]
71					
72	1	3.0	3.0	0.0	time_callback = timehistory.TimeHistory()
73					
74	1	2.0	2.0	0.0	model.fit(x_train, y_train, batch_size=self.batch_size,
75	1	1.0	1.0	0.0	epochs=self.epochs, verbose=1,
76	1	1326038.0	1326038.0	91.3	callbacks=[time_callback])
77					
78	1	5.0	5.0	0.0	self.fit_time = 0
79	2	3.0	1.5	0.0	for i in range(1, self.epochs):
80	1	1.0	1.0	0.0	self.fit_time += time_callback.times[i]
81					
82	1	0.0	0.0	0.0	self.model = model



Tensorflow - GPU

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
47					@profile
48					def run_benchmark_train(self, gpus=0):
49					
50	1	2.0	2.0	0.0	x_train, y_train = self.x_train, self.y_train
51					
52	1	5986.0	5986.0	0.5	model = Sequential()
53	1	16223.0	16223.0	1.5	model.add(Dense(512, activation='relu', input_shape=(784,)))
54	1	19485.0	19485.0	1.8	model.add(Dropout(0.2))
55	1	14582.0	14582.0	1.3	model.add(Dense(512, activation='relu'))
56	1	17566.0	17566.0	1.6	model.add(Dropout(0.2))
57	1	14073.0	14073.0	1.3	model.add(Dense(self.num_classes))
58					
59	1	3.0	3.0	0.0	if keras.backend.backend() is "tensorflow" and gpus > 1:
60					model = multi_gpu_model(model, gpus=gpus)
61					
62	1	1.0	1.0	0.0	model.compile(loss='categorical_crossentropy',
63	1	11079.0	11079.0	1.0	optimizer=RMSprop(),
64	1	28649.0	28649.0	2.6	metrics=['accuracy'])
65					
66					# create a distributed trainer for cntk
67	1	2.0	2.0	0.0	if keras.backend.backend() is "cntk" and gpus > 1:
68					start, end = cntk_gpu_mode_config(model, x_train.shape[0])
69					x_train = x_train[start: end]
70					y_train = y_train[start: end]
71					
72	1	3.0	3.0	0.0	time_callback = timehistory.TimeHistory()
73					
74	1	1.0	1.0	0.0	model.fit(x_train, y_train, batch_size=self.batch_size,
75	1	0.0	0.0	0.0	epochs=self.epochs, verbose=1,
76	1	972181.0	972181.0	88.4	callbacks=[time_callback])
77					
78	1	5.0	5.0	0.0	self.fit_time = 0
79	2	5.0	2.5	0.0	for i in range(1, self.epochs):
80	1	1.0	1.0	0.0	self.fit_time += time_callback.times[i]
81					
82	1	1.0	1.0	0.0	self.model = model

mxnet - GPU

Line #	Hits	Time	Per Hit	% Time	Line Contents
47					@profile
48					def run_benchmark_train(self, gpus=0):
49					
50	1	2.0	2.0	0.0	x_train, y_train = self.x_train, self.y_train
51					
52	1	63347.0	63347.0	16.7	model = Sequential()
53	1	8784.0	8784.0	2.3	model.add(Dense(512, activation='relu', input_shape=(784,)))
54	1	583.0	583.0	0.2	model.add(Dropout(0.2))
55	1	6107.0	6107.0	1.6	model.add(Dense(512, activation='relu'))
56	1	600.0	600.0	0.2	model.add(Dropout(0.2))
57	1	4692.0	4692.0	1.2	model.add(Dense(self.num_classes))
58					
59	1	8.0	8.0	0.0	if keras.backend.backend() is "tensorflow" and gpus > 1:
60					model = multi_gpu_model(model, gpus=gpus)
61					
62	1	1.0	1.0	0.0	model.compile(loss='categorical_crossentropy',
63	1	362.0	362.0	0.1	optimizer=RMSprop(),
64	1	4606.0	4606.0	1.2	metrics=['accuracy'])
65					
66					# create a distributed trainer for cntk
67	1	2.0	2.0	0.0	if keras.backend.backend() is "cntk" and gpus > 1:
68					start, end = cntk_gpu_mode_config(model, x_train.shape[0])
69					x_train = x_train[start: end]
70					y_train = y_train[start: end]
71					
72	1	5.0	5.0	0.0	time_callback = timehistory.TimeHistory()
73					
74	1	1.0	1.0	0.0	model.fit(x_train, y_train, batch_size=self.batch_size,
75	1	1.0	1.0	0.0	epochs=self.epochs, verbose=1,
76	1	290363.0	290363.0	76.5	callbacks=[time_callback])
77					
78	1	5.0	5.0	0.0	self.fit_time = 0
79	2	3.0	1.5	0.0	for i in range(1, self.epochs):
80	1	1.0	1.0	0.0	self.fit_time += time_callback.times[i]
81					
82	1	0.0	0.0	0.0	self.model = model

nvprof



Thank you